



:[28] In

```
import numpy as np
iris_data = np.genfromtxt('F:\\mydataset.data',
delimiter=",", skip_header=1)
print(iris_data)
```

```
[nan 0.4 1.7 3.9 5.4]
[nan 0.3 1.4 3.4 4.6]
[nan 0.2 1.5 3.4 .5]
[nan 0.2 1.4 2.9 4.4]
[nan 0.1 1.5 3.1 4.9]
[nan 0.2 1.5 3.7 5.4]
[nan 0.2 1.6 3.4 4.8]
[nan 0.1 1.4 .3 4.8]
[nan 0.1 1.1 .3 4.3]
[nan 0.2 1.2 .4 5.8]
[nan 0.4 1.5 4.4 5.7]
[nan 0.4 1.3 3.9 5.4]
[nan 0.3 1.4 3.5 5.1]
[nan 0.3 1.7 3.8 5.7]
[nan 0.3 1.5 3.8 5.1]
[nan 0.2 1.7 3.4 5.4]
[nan 0.4 1.5 3.7 5.1]
[nan 0.2 .1 3.6 4.6]
[nan 0.5 1.7 3.3 5.1]
[nan 0.2 1.9 3.4 4.8]
```



:[30] In

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(color_codes=True)
```



:[34] In


```
x = np.linspace(0, 10, 100)

fig = plt.figure()
plt.plot(x, np.sin(x), '-')
plt.plot(x, np.cos(x), '-')

```

Out[34]:

[<matplotlib.lines.Line2D at 0x28660395fc8>]



:[43] In

```
from sklearn import datasets
iris = datasets.load_iris()
print(iris.filename)
df=pd.read_csv(mydataset.filename)
df.head(5)

```

E:\Users\HP\Anaconda3\lib\site-packages\sklearn\datasets\data\iris.csv

Out[43]:

virginica	versicolor	setosa	4	150
0	0.2	1.4	3.5	5.1
0	0.2	1.4	3.0	4.9
0	0.2	1.3	3.2	4.7
0	0.2	1.5	3.1	4.6
0	0.2	1.4	3.6	5.0

```
: [41] In
```

```
<'class 'pandas.core.frame.DataFrame>
RangeIndex: 150 entries, 0 to 149
:(Data columns (total 5 columns)
non-null float64 150          150
non-null float64 150          4
setosa           150 non-null float64
versicolor      150 non-null float64
virginica        150 non-null int64
(dtypes: float64(4), int64(1)
memory usage: 6.0 KB
```

: [44] In

```
df.describe()
```

Out[44]:

virginica	versicolor	setosa	4	150	
150.000000	150.000000	150.000000	150.000000	150.000000	count
1.000000	1.199333	3.758000	3.057333	5.843333	mean
0.819232	0.762238	1.765298	0.435866	0.828066	std
0.000000	0.100000	1.000000	2.000000	4.300000	min
0.000000	0.300000	1.600000	2.800000	5.100000	25%
1.000000	1.300000	4.350000	3.000000	5.800000	50%
2.000000	1.800000	5.100000	3.300000	6.400000	75%
2.000000	2.500000	6.900000	4.400000	7.900000	max

⌂

: [47] In

```
import pandas as pd

df = pd.read_csv('F:\\mydataset.data')
df.describe()
```

Out[47]:

	0.2	1.4	3.5	5.1	
149.000000	149.000000	149.000000	149.000000	149.000000	count
1.205369	3.774497	3.051007	5.848322	1.205369	mean
0.761292	1.759651	0.433499	0.828594	0.761292	std
0.100000	1.000000	2.000000	4.300000	0.100000	min
0.300000	1.600000	2.800000	5.100000	0.300000	25%
1.300000	4.400000	3.000000	5.800000	1.300000	50%
1.800000	5.100000	3.300000	6.400000	1.800000	75%
2.500000	6.900000	4.400000	7.900000	2.500000	max

⌂

: [49] In

```
print(df.describe())
```

	0.2	1.4	3.5	5.1	
count	149.000000	149.000000	149.000000	149.000000	149.000000
mean	5.848322	3.051007	3.774497	1.205369	1.205369
std	0.828594	0.433499	1.759651	0.761292	0.761292
min	4.300000	2.000000	1.000000	0.100000	0.100000
0.300000	1.600000	2.800000	5.100000	25%	25%
1.300000	4.400000	3.000000	5.800000	50%	50%
1.800000	5.100000	3.300000	6.400000	75%	75%
max	7.900000	4.400000	6.900000	2.500000	2.500000

⌂

: [50] In

```
import pandas as pd
from sklearn import datasets
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
sns.set(color_codes=True)
```

⏮

:[54] In

```
dataset=mydataset.data()

data=pd.DataFrame(dataset['data'], columns=['petal length','petal width','sepal length','se
data['species']=dataset['target']
data['species']=data['species'].apply(lambda x: dataset['target_names'][x])
```

```
-----
(Traceback (most recent call last
<ipython-input-54-6ebe9b8539fe> in <module>
()dataset=mydataset.data 1 <----
2
data=pd.DataFrame(dataset['data'], columns=['petal length','petal wi 3
(['dth','sepal length','sepal width
['data['species']=dataset['target 4
data['species']=data['species'].apply(lambda x: dataset['target_name 5
([s']][x
```

TypeError: 'numpy.ndarray' object is not callable

⏮

:[52] In

```
data.head(10)
```

Out[52]:

species	sepal width	sepal length	petal width	petal length	
setosa	0.2	1.4	3.5	5.1	0
setosa	0.2	1.4	3.0	4.9	1
setosa	0.2	1.3	3.2	4.7	2
setosa	0.2	1.5	3.1	4.6	3
setosa	0.2	1.4	3.6	5.0	4
setosa	0.4	1.7	3.9	5.4	5
setosa	0.3	1.4	3.4	4.6	6
setosa	0.2	1.5	3.4	5.0	7
setosa	0.2	1.4	2.9	4.4	8
setosa	0.1	1.5	3.1	4.9	9

⏮

:[55] In

```
data.isnull().sum() #there is no missing value.
```

Out[55]:

```
petal length    0
petal width     0
sepal length    0
sepal width     0
species         0
dtype: int64
```



:[56] In

```
modData = data.append({'petal length' : np.nan , 'petal width' : 3.6, 'sepal length' : 0,
                      'sepal width': 0.2, 'species' : 'setosa'} ,ignore_index=True)
modData.describe()
```

Out[56]:

sepal width	sepal length	petal width	petal length	
151.000000	151.000000	151.000000	150.000000	count
1.192715	3.733113	3.060927	5.843333	mean
0.764033	1.785785	0.436650	0.828066	std
0.100000	0.000000	2.000000	4.300000	min
0.300000	1.550000	2.800000	5.100000	25%
1.300000	4.300000	3.000000	5.800000	50%
1.800000	5.100000	3.350000	6.400000	75%
2.500000	6.900000	4.400000	7.900000	max



:[57] In

```
print('Columns with missing values')
print(modData.isnull().sum())
print('\n Columns with zero values')
print((modData[['petal length','petal width','sepal length','sepal width','species']]==0).sum())
```

Columns with missing values

```
petal length    1
petal width     0
sepal length    0
sepal width     0
species         0
dtype: int64
```

Columns with zero values

```
petal length    0
petal width     0
sepal length    1
sepal width     0
species         0
dtype: int64
```



:[58] In

```
modData.fillna(modData.mean(), inplace=True)
print(modData.isnull().sum())
```

```
petal length    0
petal width     0
sepal length    0
sepal width     0
species         0
dtype: int64
```

:[62] In

```
url = "https://archive.ics.uci.edu/ml/machine-learning-databases/iris/iris.data"
PCA_df = pd.read_csv(url, names=['sepal length', 'sepal width', 'petal length', 'petal width'],
```

:[63] In

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
features = ['sepal length', 'sepal width', 'petal length', 'petal width']
# Separating out the features
x = PCA_df.loc[:, features].values
# Separating out the target
y = PCA_df.loc[:, ['target']].values
# Standardizing the features
x = StandardScaler().fit_transform(x)
print (x)
```

```
[9.00681170e-01  1.03205722e+00 -1.34127240e+00 -1.31297673e+00-]
[1.14301691e+00 -1.24957601e-01 -1.34127240e+00 -1.31297673e+00-]
[1.38535265e+00  3.37848329e-01 -1.39813811e+00 -1.31297673e+00-]
[1.50652052e+00  1.06445364e-01 -1.28440670e+00 -1.31297673e+00-]
[1.02184904e+00  1.26346019e+00 -1.34127240e+00 -1.31297673e+00-]
[5.37177559e-01  1.95766909e+00 -1.17067529e+00 -1.05003079e+00-]
[1.50652052e+00  8.00654259e-01 -1.34127240e+00 -1.18150376e+00-]
[1.02184904e+00  8.00654259e-01 -1.28440670e+00 -1.31297673e+00-]
[1.74885626e+00 -3.56360566e-01 -1.34127240e+00 -1.31297673e+00-]
[1.14301691e+00  1.06445364e-01 -1.28440670e+00 -1.44444970e+00-]
[5.37177559e-01  1.49486315e+00 -1.28440670e+00 -1.31297673e+00-]
[1.26418478e+00  8.00654259e-01 -1.22754100e+00 -1.31297673e+00-]
[1.26418478e+00 -1.24957601e-01 -1.34127240e+00 -1.44444970e+00-]
[1.87002413e+00 -1.24957601e-01 -1.51186952e+00 -1.44444970e+00-]
[5.25060772e-02  2.18907205e+00 -1.45500381e+00 -1.31297673e+00-]
[1.73673948e-01  3.11468391e+00 -1.28440670e+00 -1.05003079e+00-]
[5.37177559e-01  1.95766909e+00 -1.39813811e+00 -1.05003079e+00-]
[9.00681170e-01  1.03205722e+00 -1.34127240e+00 -1.18150376e+00-]
[1.73673948e-01  1.72626612e+00 -1.17067529e+00 -1.18150376e+00-]
```

:[64] In

```
from sklearn.decomposition import PCA
pca = PCA(n_components = 2)
principalComponents = pca.fit_transform(x)
principalDf = pd.DataFrame(data = principalComponents
                           , columns = ['principal component 1', 'principal component 2'])
```


:[65] In

```
finalDf = pd.concat([principalDf, PCA_df[['target']]], axis = 1)
```

:[66] In


```
import matplotlib.pyplot as plt
%matplotlib inline
#sns.set(color_codes=True)

fig = plt.figure(figsize = (8,8))
ax = fig.add_subplot(1,1,1)
ax.set_xlabel('Principal Component 1', fontsize = 15)
ax.set_ylabel('Principal Component 2', fontsize = 15)
ax.set_title('2 component PCA', fontsize = 20)
targets = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
colors = ['r', 'g', 'b']
for target, color in zip(targets,colors):
    indicesToKeep = finalDf['target'] == target
    ax.scatter(finalDf.loc[indicesToKeep, 'principal component 1']
              , finalDf.loc[indicesToKeep, 'principal component 2']
              , c = color
              , s = 50)
ax.legend(targets)
ax.grid()
```



:[73] In

```
Error = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters = i).fit(x)
    kmeans.fit(x)
    Error.append(kmeans.inertia_)
plt.plot(range(1, 11), Error)
plt.title('Elbow method')
plt.xlabel('No of clusters')
plt.ylabel('Error')
plt.show()
```




:[74] In

```
plt.scatter(x[:,0], x[:,1], c=y_kmeans5, cmap='rainbow')
```

Out[74]:

<matplotlib.collections.PathCollection at 0x28660f9f088>



K

:[75] In

```
pip install mlxtend
```

```
Requirement already satisfied: mlxtend in e:\users\hp\anaconda3\lib\site-packages (0.17.2)
Requirement already satisfied: pandas>=0.24.2 in e:\users\hp\anaconda3\lib\site-packages (from mlxtend) (0.25.1)
Requirement already satisfied: matplotlib>=3.0.0 in e:\users\hp\anaconda3\lib\site-packages (from mlxtend) (3.1.1)
Requirement already satisfied: joblib>=0.13.2 in e:\users\hp\anaconda3\lib\site-packages (from mlxtend) (0.13.2)
Requirement already satisfied: numpy>=1.16.2 in e:\users\hp\anaconda3\lib\site-packages (from mlxtend) (1.16.5)
Requirement already satisfied: setuptools in e:\users\hp\anaconda3\lib\site-packages (from mlxtend) (41.4.0)
Requirement already satisfied: scikit-learn>=0.20.3 in e:\users\hp\anaconda3\lib\site-packages (from mlxtend) (0.21.3)
Requirement already satisfied: scipy>=1.2.1 in e:\users\hp\anaconda3\lib\site-packages (from mlxtend) (1.3.1)
Requirement already satisfied: python-dateutil>=2.6.1 in e:\users\hp\anaconda3\lib\site-packages (from pandas>=0.24.2->mlxtend) (2.8.0)
Requirement already satisfied: pytz>=2017.2 in e:\users\hp\anaconda3\lib\site-packages (from pandas>=0.24.2->mlxtend) (2019.3)
Requirement already satisfied: cyclor>=0.10 in e:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in e:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (1.1.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in e:\users\hp\anaconda3\lib\site-packages (from matplotlib>=3.0.0->mlxtend) (2.4.2)
Requirement already satisfied: six>=1.5 in e:\users\hp\anaconda3\lib\site-packages (from python-dateutil>=2.6.1->pandas>=0.24.2->mlxtend) (1.12.0)
Note: you may need to restart the kernel to use updated packages
```

K

:[76] In

```
from mlxtend.frequent_patterns import apriori
```

⌵

: [77] In

dataset

Out[77]:

,[data': array([[5.1, 3.5, 1.4, 0.2']
,[0.2 ,1.4 , .3 ,4.9]
,[0.2 ,1.3 ,3.2 ,4.7]
,[0.2 ,1.5 ,3.1 ,4.6]
,[0.2 ,1.4 ,3.6 , .5]
,[0.4 ,1.7 ,3.9 ,5.4]
,[0.3 ,1.4 ,3.4 ,4.6]
,[0.2 ,1.5 ,3.4 , .5]
,[0.2 ,1.4 ,2.9 ,4.4]
,[0.1 ,1.5 ,3.1 ,4.9]
,[0.2 ,1.5 ,3.7 ,5.4]
,[0.2 ,1.6 ,3.4 ,4.8]
,[0.1 ,1.4 , .3 ,4.8]
,[0.1 ,1.1 , .3 ,4.3]
,[0.2 ,1.2 , .4 ,5.8]
,[0.4 ,1.5 ,4.4 ,5.7]
,[0.4 ,1.3 ,3.9 ,5.4]

⌵

: [78] In

```
import pandas as pd
from mlxtend.preprocessing import TransactionEncoder

TranEncod = TransactionEncoder()
te_ary = TranEncod.fit(dataset).transform(dataset)
df = pd.DataFrame(te_ary, columns=TranEncod.columns_)
```

⌵

: [79] In

df

Out[79]:

	m	l	i	g	f	e	d	a	_	S	R	E	D	C	
False	False	False	False	False	False	True	True	False	False	False	False	False	False	False	0
False	False	False	True	False	True	False	True	False	False	False	False	False	False	False	1
True	False	False	True	False	True	False	True	True	False	False	False	False	False	False	2
False	False	False	False	False	False	False	False	False	True	True	True	True	True	True	3
True	False	False	False	True	True	False	True	True	False	False	False	False	False	False	4
True	True	True	False	True	True	False	True	False	False	False	False	False	False	False	5

⏮

:[80] In

```
apriori(df, min_support=0.6)
```

Out[80]:

itemsets	support	
(6)	0.833333	0
(8)	0.666667	1
(17)	0.666667	2
(6 ,8)	0.666667	3
(6 ,17)	0.666667	4

⏮

:[81] In

```
apriori(df, min_support=0.6, use_colnames=True)
```

Out[81]:

itemsets	support	
(a)	0.833333	0
(e)	0.666667	1
(t)	0.666667	2
(a, e)	0.666667	3
(a, t)	0.666667	4

⏮

:[82] In

apriori(df, min_support=0.3, use_colnames=True)

Out[82]:

itemsets	support	
(_)	0.333333	0
(a)	0.833333	1
(e)	0.666667	2
(f)	0.333333	3
(g)	0.333333	4
...
(a, r, t, e, s, n, _)	0.333333	282
(m, a, r, t, s, n, _)	0.333333	283
(m, r, t, e, s, n, _)	0.333333	284
(m, a, r, t, e, s, n)	0.333333	285
(m, a, r, t, e, s, n, _)	0.333333	286
rows × 2 columns 287		

⌕

: [83] In

```
frequent_itemsets = apriori(df, min_support=0.5, use_colnames=True)
frequent_itemsets['length'] = frequent_itemsets['itemsets'].apply(lambda x: len(x))
frequent_itemsets
```

Out[83]:

length	itemsets	support	
1	(a)	0.833333	0
1	(e)	0.666667	1
1	(m)	0.500000	2
1	(n)	0.500000	3
1	(r)	0.500000	4
1	(t)	0.666667	5
2	(a, e)	0.666667	6
2	(m, a)	0.500000	7
2	(a, n)	0.500000	8
2	(a, r)	0.500000	9
2	(a, t)	0.666667	10
2	(m, e)	0.500000	11
2	(e, n)	0.500000	12
2	(e, r)	0.500000	13
2	(e, t)	0.500000	14
2	(m, n)	0.500000	15
2	(r, t)	0.500000	16
3	(m, a, e)	0.500000	17
3	(a, e, n)	0.500000	18
3	(a, e, r)	0.500000	19
3	(a, e, t)	0.500000	20
3	(m, a, n)	0.500000	21
3	(a, r, t)	0.500000	22
3	(m, e, n)	0.500000	23
3	(e, r, t)	0.500000	24
4	(m, a, e, n)	0.500000	25
4	(a, e, r, t)	0.500000	26

⏮

:[84] In

```
frequent_itemsets[ (frequent_itemsets['length'] == 2) &
                    (frequent_itemsets['support'] >= 0.5) ]
```

Out[84]:

length	itemsets	support	
2	(a, e)	0.666667	6
2	(m, a)	0.500000	7
2	(a, n)	0.500000	8
2	(a, r)	0.500000	9
2	(a, t)	0.666667	10
2	(m, e)	0.500000	11
2	(e, n)	0.500000	12
2	(e, r)	0.500000	13
2	(e, t)	0.500000	14
2	(m, n)	0.500000	15
2	(r, t)	0.500000	16

⏮

:[85] In

```
frequent_itemsets[ frequent_itemsets['itemsets'] == {'Diaper', 'Drink'} ]
```

Out[85]:

length	itemsets	support
--------	----------	---------

⏮

:[86] In

```
pip install Graphviz
```

Requirement already satisfied: Graphviz in e:\users\hp\anaconda3\lib\site-packages (0.13.2)
 .Note: you may need to restart the kernel to use updated packages

⏮

:[87] In

```
import pandas as pd
import graphviz
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
# Setting random seed.
seed = 10
```


⏮

:[88] In

```

from sklearn import datasets
iris= datasets.load_iris()
df=pd.read_csv('F:\mydataset.data', delimiter=',', header= 0, names= ['sepal length (cm)',
df.head()

```

Out[88]:

Variety	(petal width (cm	(petal length (cm	(sepal width (cm	(sepal length (cm	
Iris-setosa	0.2	1.4	3.0	4.9	0
Iris-setosa	0.2	1.3	3.2	4.7	1
Iris-setosa	0.2	1.5	3.1	4.6	2
Iris-setosa	0.2	1.4	3.6	5.0	3
Iris-setosa	0.4	1.7	3.9	5.4	4

⏮

:[89] In

```

# Creating a LabelEncoder and fitting it to the dataset labels.
le = LabelEncoder()
le.fit(df['Variety'].values)
# Converting dataset str labels to int labels.
y = le.transform(df['Variety'].values)
# Extracting the instances data.
X = df.drop('Variety', axis=1).values
# Splitting into train and test sets.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.34, stratify=y, random

```

⏮

:[90] In

```

# Creating a DecisionTreeClassifier.
# The criterion parameter indicates the measure used (possible values: 'gini' for the Gini
# 'entropy' for the information gain).
# The min_samples_leaf parameter indicates the minimum of objects required at a leaf node.
# The min_samples_split parameter indicates the minimum number of objects required to split
# The max_depth parameter controls the maximum tree depth. Setting this parameter to None w
# tree until all leaves are pure or until all leaves contain less than min_samples_split sa
tree = DecisionTreeClassifier(criterion='gini',
                             min_samples_leaf=5,
                             min_samples_split=5,
                             max_depth=None,
                             random_state=seed)
tree.fit(X_train, y_train)
y_pred = tree.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print('DecisionTreeClassifier accuracy score: {}'.format(accuracy))

```

DecisionTreeClassifier accuracy score: 0.9607843137254902

K

:[91] In


```

from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt

print('Confusion Matrix is')
print(confusion_matrix(y_test, y_pred))
cm=confusion_matrix(y_test, y_pred)
plt.matshow(cm)
plt.show()

```

Confusion Matrix is
 [0 0 17]
 [1 16 0]
 [16 1 0]



K

:[98] In

```

# Assigning features and Label variables
weather=['Rainy','Rainy','Overcast','Sunny','Sunny','Sunny',
         'Overcast','Rainy','Rainy','Sunny','Rainy','Overcast',
         'Overcast','Sunny']

temp=['Hot','Hot','Hot','Mild','Cool','Cool','Cool',
      'Mild','Cool','Mild','Mild','Mild','Hot','Mild']

play=['No','No','Yes','Yes','Yes','No','Yes','No','Yes','Yes','Yes','Yes','Yes','Yes','No']

```

⏮

: [99] In

```
import pandas as pd
data= {'weather': ['Rainy','Rainy','Overcast','Sunny','Sunny','Sunny',
                  'Overcast','Rainy','Rainy','Sunny','Rainy','Overcast',
                  'Overcast','Sunny'],
       'temp': ['Hot','Hot','Hot','Mild','Cool','Cool','Cool',
               'Mild','Cool','Mild','Mild','Mild','Hot','Mild'],
       'play': ['No','No','Yes','Yes','Yes','No','Yes','No','Yes','Yes','Yes','Yes','Yes','N

data= pd.DataFrame(data)
data
```

Out[99]:

play	temp	weather	
No	Hot	Rainy	0
No	Hot	Rainy	1
Yes	Hot	Overcast	2
Yes	Mild	Sunny	3
Yes	Cool	Sunny	4
No	Cool	Sunny	5
Yes	Cool	Overcast	6
No	Mild	Rainy	7
Yes	Cool	Rainy	8
Yes	Mild	Sunny	9
Yes	Mild	Rainy	10
Yes	Mild	Overcast	11
Yes	Hot	Overcast	12
No	Mild	Sunny	13

⏮

: [105] In

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
```

⏮

: [106] In

```
from sklearn import datasets

dataset= datasets.load_iris()
df= pd.DataFrame(dataset['data'], columns=['sepal length (cm)', 'sepal width (cm)', 'petal l
```



:[107] In

```
X= df['sepal length (cm)']
Y= df['petal length (cm)']
Slic_df = pd.DataFrame({'sepal length': X, 'petal length': Y})
print(Slic_df)
```


sepal length	petal length	
1.4	5.1	0
1.4	4.9	1
1.3	4.7	2
1.5	4.6	3
1.4	5.0	4
...
5.2	6.7	145
5.0	6.3	146
5.2	6.5	147
5.4	6.2	148
5.1	5.9	149

[rows x 2 columns 150]



:[108] In

```
plt.scatter(Slic_df[['sepal length']], Slic_df[['petal length']], color = "r", marker = "x")
plt.xlabel('Sepal length')
plt.ylabel('Petal length')
plt.title('Scatter Plot')
plt.show()
```



:[109] In

```
from sklearn.linear_model import LinearRegression

#define the classifier
classifier = LinearRegression()

#train the classifier
model = classifier.fit(Slic_df[['sepal length']], Slic_df[['petal length']])
```

:[110] In

```

#use the trained classifier to make prediction
y_pred = classifier.predict(Slic_df[['sepal length']])
print(y_pred)

#print coefficient (a in y=ax+b) and intercept (the constant, b in y=ax+b)
print('Coefficients: \n', classifier.coef_)
print('Intercept: \n', classifier.intercept_)

```

```

[2.37656482]
[2.00487822]
[1.63319163]
[1.44734833]
[2.19072152]
[2.93409471]
[1.44734833]
[2.19072152]
[1.07566173]
[2.00487822]
[2.93409471]
[1.81903493]
[1.81903493]
[0.88981844]
[ 3.6774679]
[3.49162461]
[2.93409471]
[2.37656482]
[3.49162461]

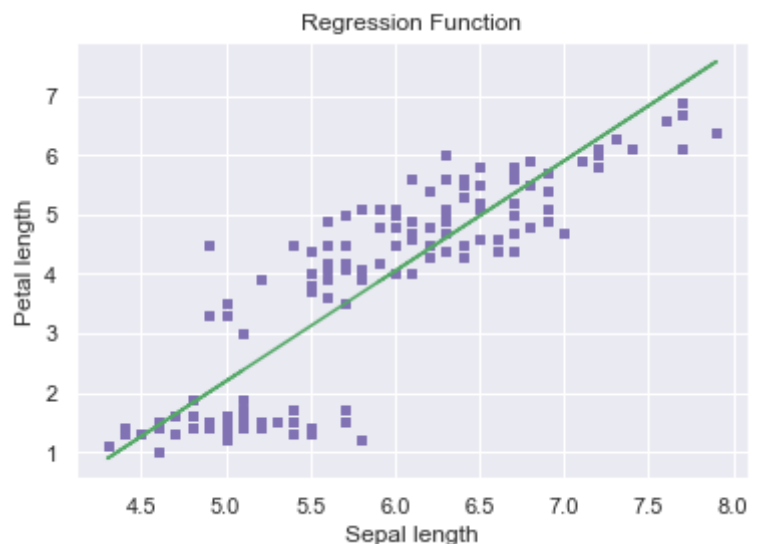
```

:[111] In

```

#visualize original data points
plt.scatter(Slic_df[['sepal length']], Slic_df[['petal length']], color = "m", marker = "s")
#visualize regression function
plt.plot(Slic_df['sepal length'], y_pred, color = "g")
plt.xlabel('Sepal length')
plt.ylabel('Petal length')
plt.title('Regression Function')
plt.show()

```



:[112] In

```

df2 = pd.DataFrame({'col1': [19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31],
                    'col2': [36, 70, 48, 119, 51, 205, 133, 112, 92, 99, 96, 154, 110]})


#your turn - solution

#define the classifier
classifier = LinearRegression()
#train the classifier
classifier.fit(df2[['col1']], df2[['col2']])

#use the classifier to make prediction
y_predict = classifier.predict(df2[['col1']])

#visualize data points
plt.scatter(df2[['col1']], df2[['col2']], color = "m", marker = "o", s = 30)
#visualize regression function
plt.plot(df2[['col1']], y_predict, color = "g")
plt.xlabel('x')
plt.ylabel('y')
plt.title('title')
plt.show()
#print coefficient (a in y=ax+b) and intercept (the constant, b in y=ax+b)
print('Coefficients: \n', classifier.coef_)
print('Intercept: \n', classifier.intercept_)

```



```

:Coefficients
[[5.41208791]]
:Intercept
[33.37912088-]

```

:[113] In

```

Day_18 = classifier.intercept_ + classifier.coef_ * 18
print(Day_18)

```

```

[[64.03846154]]

```

K

:[114] In

```
from sklearn.metrics import mean_squared_error, mean_absolute_error


print("Mean squared error:")
print(mean_squared_error(df2[['col2']], y_predict))
```

```
:Mean squared error
1572.1551141166526
```

K

:[115] In

```
plt.scatter(y_predict, (df2[['col2']] - y_predict) ** 2, color = "blue", s = 10,)
plt.title("Squared errors")
## plotting line for zero error
plt.hlines(y = 0, xmin = 0, xmax = 140, linewidth = 2)
plt.xlabel('predicted values')
plt.ylabel('squared error')
plt.show()
```



:[116] In

```

#Your turn
print("Mean absolute error: ")
print(mean_absolute_error(df2['col2'], y_predict))


#plot
plt.scatter(y_predict, (df2[['col2']] - y_predict), color = "blue", s = 10,)
plt.title("Absolute Errors")
plt.hlines(y = 0, xmin = 0, xmax = 140, linewidth = 2)
plt.xlabel('predicted values')
plt.ylabel('absolute error')
plt.show()

```

```

:Mean absolute error
31.163144547759927

```



:[117] In

```

df['Species'] = dataset['target']
df['Species'] = df['Species'].apply(lambda x: dataset['target_names'][x])
df['Species']

```

Out[117]:

```

setosa      0
setosa      1
setosa      2
setosa      3
setosa      4
...
virginica   145
virginica   146
virginica   147
virginica   148
virginica   149
Name: Species, Length: 150, dtype: object

```




:[118] In

```
from sklearn.linear_model import LogisticRegression

X = Slic_df[['sepal length']] #we only use the first feature as descriptive feature
y = df['Species'] #use the species as target feature

#defining and training the Logistic regression model
log_regression = LogisticRegression(solver = 'liblinear', multi_class = 'ovr')
log_regression.fit(X, y)

pred = log_regression.predict(X)

print('Score: \n', log_regression.score(X, y)) #the mean accuracy on the given test data and
print('Coefficients: \n', log_regression.coef_)
print('Intercept: \n', log_regression.intercept_)
```

```
      :Score
0.6933333333333334
      :Coefficients
[0.86959145-]
[0.01223362 ]
[[0.57972675 ]
      :Intercept
[ 3.9921824- 0.74244291- 4.16186636 ]
```

⏮

:[119] In

```

#we use all descriptive features, but not the target feature
X = df.iloc[:, 0:4]
y=df['Species']
print('X: \n', X.head(), '\n')
print('y: \n', y.head(), '\n')

from sklearn.linear_model import LogisticRegression
classifier3 = LogisticRegression( solver = 'liblinear', multi_class = 'ovr')
classifier3.fit(X, y)

y_pred = classifier3.predict(X)

print('Score: \n', classifier3.score(X, y))
print('Coefficients: \n', classifier3.coef_)
print('Intercept: \n', classifier3.intercept_)

```

```

:X
(sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0.2                1.4                3.5                5.1                0
0.2                1.4                3.0                4.9                1
0.2                1.3                3.2                4.7                2
0.2                1.5                3.1                4.6                3
0.2                1.4                3.6                5.0                4

:y
setosa    0
setosa    1
setosa    2
setosa    3
setosa    4
Name: Species, dtype: object

:Score
0.96
:Coefficients
[1.02103509- 2.26003266- 1.46416217  0.41021713 ]
[1.40617325- 0.5758173  1.61211605- 0.4275087 ]
[[2.55537041  2.47096755  1.53427768- 1.70751526-]
:Intercept
[1.21470917- 1.09392467  0.26421853 ]

```

⏮

:[] In